

ADITYA

Generative AI Engineer

aditya78332412@gmail.com | [+91-7833934588](tel:+91-7833934588) | Noida, India | linkedin.com/in/adityasanjay1999
github.com/aditya2425 | [Portfolio](#)

PROFESSIONAL SUMMARY

Generative AI Engineer with 3+ years at Accenture, building production-grade AI systems for enterprise clients across healthcare, aviation, and entertainment. Hands-on with RAG pipelines, multi-agent orchestration (Semantic Kernel, LangChain, CrewAI), LLM fine-tuning (QLoRA, RLHF), LLM evaluation (RAGAS, BERTScore), and safety guardrails (PII detection, content filtering). Shipped end-to-end GenAI applications independently covering the full lifecycle from data preparation to production monitoring.

TECHNICAL SKILLS

GenAI & LLMs: OpenAI GPT-4o, Anthropic Claude, Mistral 7B, Hugging Face Transformers, Prompt Engineering (CoT, Few-Shot, ReAct), Retrieval-Augmented Generation (RAG), Multi-Agent Systems, Agentic AI, LLM Fine-Tuning (QLoRA, LoRA, PEFT, RLHF, SFTTrainer), Multimodal RAG

Agentic AI & Orchestration: LangChain, LangGraph, CrewAI, Semantic Kernel, Azure AI Foundry, Multi-Agent Orchestration, Tool Calling, Function Calling, MCP (Model Context Protocol), ReAct Pattern, Human-in-the-Loop

RAG & Vector Search: Azure AI Search, Azure Cognitive Search, Pinecone, ChromaDB, BM25, Hybrid Search (Keyword + Vector), Cohere Reranking, Vector Embeddings, Semantic Search, Reciprocal Rank Fusion (RRF)

AI Safety, Evaluation & MLOps: LLM Evaluation (RAGAS-style, BERTScore, BLEU, ROUGE, LLM-as-Judge), Guardrails, Content Filtering, PII Detection, Prompt Injection Prevention, Model Monitoring, Drift Detection, A/B Testing, GitHub Actions CI/CD

Cloud & Infrastructure: Microsoft Azure (AI Foundry, Functions, App Service, Event Hub, Logic Apps, Blob Storage, Key Vault), Docker, Docker Compose, FastAPI, Health Checks, Security Scanning, Azure DevOps

Backend & Data: Python, SQL, Flask, Cosmos DB, Azure SQL, SQLite, REST APIs, Streamlit, Gradio, Pydantic, Git, pytest, Event-Driven Architecture, Microservices

PROFESSIONAL EXPERIENCE

Software Engineer - Generative AI

Jul 2024 – Present | Gurugram, HR

Accenture Solutions Private Limited

- Designed and deployed a multi-agent AI system for a global airline technology provider using Azure AI Foundry and Semantic Kernel — automating failure triage, test generation, and regression scoping, reducing resolution time from 1–2 days to under 10 minutes.
- Built Hybrid RAG pipelines with deduplication, human-in-the-loop workflows, and a QA chatbot achieving 90%+ accuracy with sub-60s response.
- Engineered agentic test case generation with regression impact prediction, risk scoring (1–5), and data ingestion pipelines into Cosmos DB with daily incremental refresh

Associate Software Engineer

Jan 2023 – Jun 2024 | Gurugram, HR

Accenture Solutions Private Limited

- Architected a real-time AI assistant for a Fortune 10 healthcare company processing thousands of live call utterances daily using Flask, Azure OpenAI, and Cognitive Search with sub-second latency.
- Built multiple production APIs covering NLP identity verification, real-time intent detection (CLU + GPT fallback), confidence-scored auto-suggestions, and post-call summarization.
- Designed event-driven pipeline (Event Hub → Functions → CosmosDB → Flask → React) with multi-threaded execution and feedback loops across multiple repositories.
- Delivered a large-scale GenAI application for a Fortune 100 entertainment company using GPT-4o — automating document generation, competency mapping, and interview guides, reducing manual effort by ~70–80%.
- Engineered specialized prompt workflows with context injection, inclusive language detection across multiple bias categories, and document validation against quality criteria.
- Migrated large-scale historical records to Azure Blob Storage with automated refresh. Delivered within 4 months with a hybrid US-India team.

KEY PROJECTS

DocuMind — Production RAG Pipeline

- Python, OpenAI, ChromaDB, Pinecone, FastAPI, Docker
- Built production RAG pipeline with 3 chunking strategies, hybrid search (dense + BM25), Cohere reranking, and multi-LLM cited answer generation
 - Developed RAGAS-style evaluation with 6 preset experiment configurations and interactive results dashboard

AgentForge — Multi-Agent Research System

- Python, LangGraph, CrewAI, Tavily, ChromaDB, Streamlit
- Designed 5-agent research pipeline with reflection loops in both LangGraph and CrewAI to benchmark framework trade-offs
 - Integrated guardrails for prompt injection, PII masking, hallucination scoring, and loop prevention

ModelForge — LLM Fine-Tuning Pipeline

- Python, HuggingFace, QLoRA, PEFT, BERTScore, Streamlit
- Fine-tuned Mistral 7B using QLoRA on 500+ code review examples with SHA256 and MinHash deduplication
 - Built evaluation dashboard comparing 4 models across BLEU, ROUGE, BERTScore, and LLM-as-Judge with cost breakeven analysis

GuardianAI — LLM Safety & Observability Platform

- Python, SQLite, Streamlit, Plotly
- Built decorator-based LLM tracing capturing latency, tokens, cost, and 4 quality metrics across 9 model pricing tiers
 - Implemented drift detection, configurable alert rules, LRU cache, and A/B testing framework

ShipAI — ML Deployment & DevOps Pipeline

- Python, Docker, GitHub Actions, FastAPI, CI/CD
- Created CI/CD pipeline with GitHub Actions, multi-stage Docker builds, and security scanning for dependencies and Dockerfiles
 - Designed 3-tier health checks and deployment runbooks with step-by-step rollback procedures

EDUCATION

B.Tech in Computer Science	2018 – 2022 Noida , UP
Amity University, Noida	

ACHIEVEMENTS

Sparkling Star Award	2024
GenAI contributions, Accenture & Avanade	
ACE Award (Team)	2023
Exceptional delivery on critical projects	

CERTIFICATIONS

Azure AI Engineer (AI-102) Azure Developer (AZ-204) Azure AI Fundamentals (AI-900) AWS Cloud Practitioner Azure Fundamentals (AZ-900) MongoDB Certified
